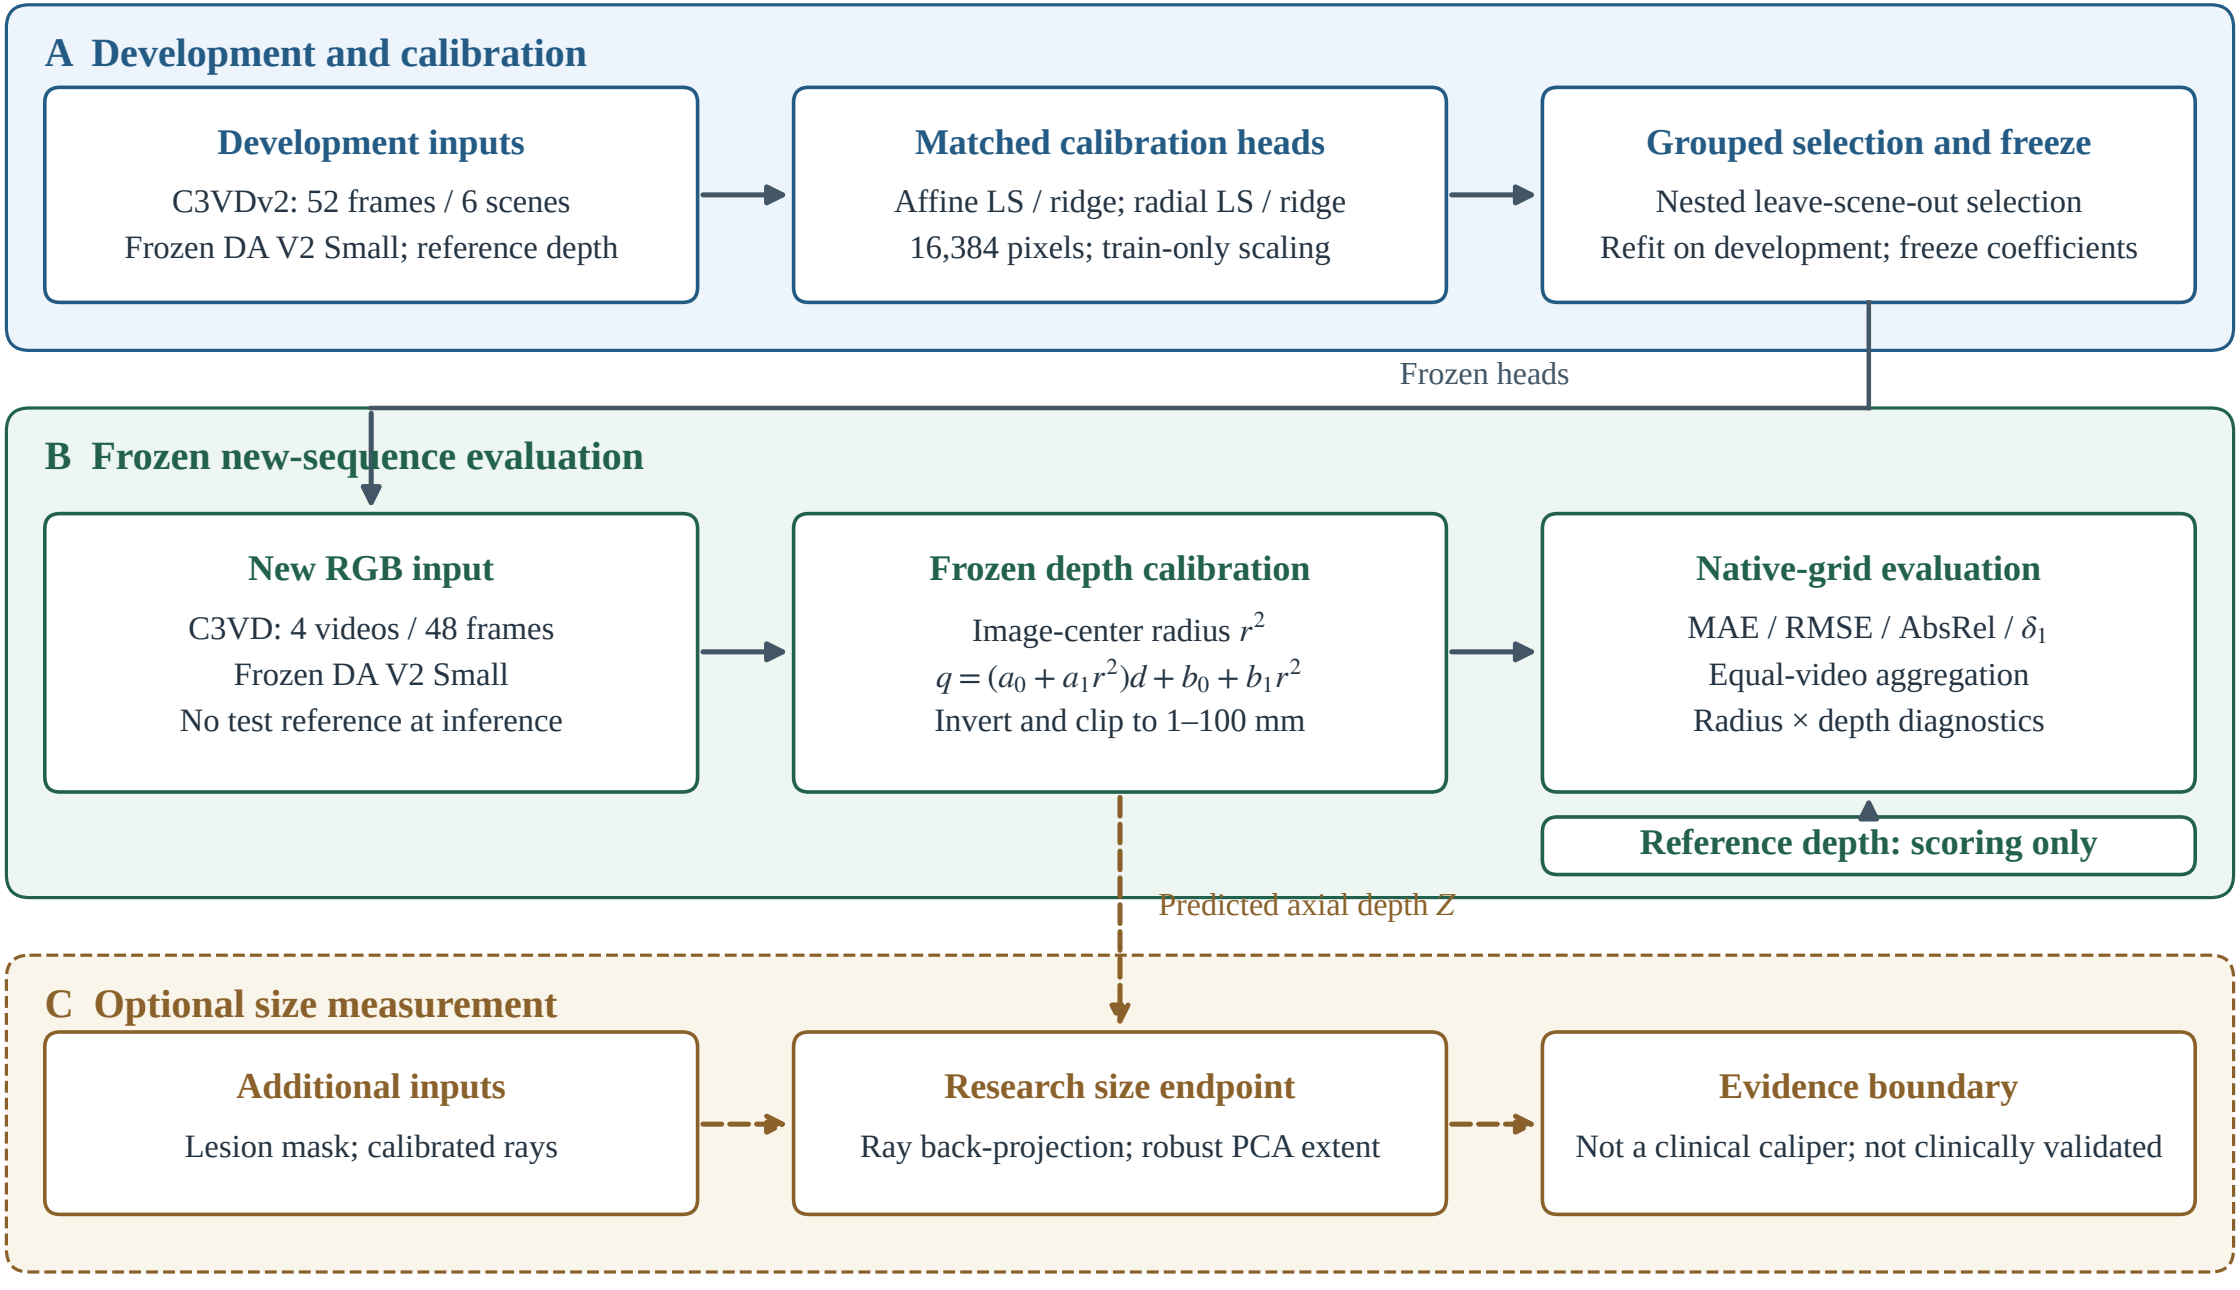
A Development and calibration
Development inputs
C3VDv2: 52 frames / 6 scenes
Frozen DA V2 Small; reference depth
Matched calibration heads
Affine LS / ridge; radial LS / ridge
16,384 pixels; train-only scaling
Grouped selection and freeze
Nested leave-scene-out selection
Refit on development; freeze coefficients
Frozen heads
B Frozen new-sequence evaluation
New RGB input
C3VD: 4 videos / 48 frames
Frozen DA V2 Small
No test reference at inference
Frozen depth calibration
Image-center radius r^2
$q = (a_0 + a_1 r^2)d + b_0 + b_1 r^2$
Invert and clip to 1–100 mm
Native-grid evaluation
MAE / RMSE / AbsRel / δ_1
Equal-video aggregation
Radius × depth diagnostics
Reference depth: scoring only
Predicted axial depth Z
C Optional size measurement
Additional inputs
Lesion mask; calibrated rays
Research size endpoint
Ray back-projection; robust PCA extent
Evidence boundary
Not a clinical caliper; not clinically validated